ORACLE SOLARIS ZFS

A BETTER SAFER WAY TO MANAGE YOUR DATA

KEY FEATURES

- No silent data corruption ever. The Oracle Solaris ZFS self healing feature automatically repairs corrupt data.
- Rock-solid data integrity. All data is protected by 256-bit checksums and data consistency is maintained at all times.
- Mind-boggling scalability: 128-bit file system, 16 billion billion times the capacity of 32- or 64-bit file systems
- Breathtaking speed: proven cutting edge technologies combine to optimize performance.
- Near-zero administration: complicated storage administration is automated and simplified, reducing administrative overhead by up to 80%
- Built-in data services: encryption, compression, deduplication, infinite snapshots, clones, replication,
- Seamless integration of data management protocols, including NFS, CIFS, and iSCSI.

Anyone who has ever lost important files, run out of space on a partition, spent weekends adding new storage to servers, tried to grow or shrink a file system, or experienced data corruption knows that there is room for improvement in file systems and volume managers. Oracle Solaris ZFS is designed from the ground up to meet the emerging needs of a general-purpose local file system that spans the desktop to the data center.

The Future-Proof File System

Oracle Solaris ZFS offers a dramatic advance in data management with an innovative approach to data integrity, near-zero administration, and a welcome integration of file system and volume management capabilities. The centerpiece of this new architecture is the concept of a virtual storage pool, which decouples the file system from physical storage in the same way that virtual memory abstracts physical memory, allowing for much more efficient and flexible use of storage devices. In Oracle Solaris ZFS, space is shared dynamically among multiple file systems from a single storage pool and is parceled out of the pool as file systems request it. Physical storage can be added to storage pools dynamically without interrupting services, providing new levels of flexibility, availability, and performance. And in terms of scalability, the theoretical limits of the 128-bit ZFS address space are truly mind-boggling: 16 exabytes (16×10 raised to the power of 18 bytes) being the maximum size of a single file.

Oracle Solaris ZFS implements an improvement on RAID-5, RAID-Z3, which uses parity, striping, and atomic operations to ensure reconstruction of corrupted data even in the face of three concurrent drive failures. It is ideally suited for managing industry standard storage servers.

Cutting-Edge Data Integrity

Oracle Solaris ZFS combines proven and cutting-edge technologies, such as copy-on-write and end-to-end checksumming. Data is always written to a new block on disk before changing the pointers to the data and committing the write. And, because the file system is always consistent, time-consuming recovery procedures like *fsck* are not required, even if the system is shut down in an unclean manner. Copy-on-write also enables administrators to take consistent backups or roll data back to a known point in time.

Oracle Solaris with Oracle Solaris ZFS is the only known operating system designed to provide end-to-end checksumming for all data. Oracle Solaris ZFS constantly reads and checks data to help ensure that it is correct, and if it detects an error in a mirrored pool, the technology can automatically repair the corrupt data. This relentless vigilance on behalf of availability protects against costly and time-consuming data loss (even previously undetectable silent data corruption).

High Performance

This radical new architecture optimizes and simplifies code paths from the application to the hardware, producing sustained throughput at exceptional speeds. New block allocation

algorithms accelerate write operations, consolidating what would traditionally be many small random writes into a single, more efficient sequential operation.

Additionally, Oracle Solaris ZFS implements intelligent prefetch, performing read-ahead for sequential data streaming, and can adapt its read behavior on the fly for more complex access patterns. To eliminate bottlenecks and increase the speed of both reads and writes, Oracle Solaris ZFS stripes data across all available storage devices, balancing I/O and maximizing throughput. And, as disks are added to the storage pool, Oracle Solaris ZFS immediately begins to allocate blocks from those devices, increasing effective bandwidth as each device is added. This means that system administrators no longer need to monitor storage devices to see if they are causing I/O bottlenecks.

Oracle Solaris ZFS is optimized to use Solid State Drives (SSDs) as intelligent cache for the most demanding workloads.

Simplified Administration

Because most file system administration tasks are painful, slow operations that are infrequently performed, they are notably prone to operator errors with the potential to destroy a great amount of data very quickly. Oracle Solaris ZFS helps alleviate this problem by automating both common and less frequent administrative tasks.

In Oracle Solaris 11, ZFS is the default file system and this file system technology underpins installation, patching and upgrade/rollback. Users can easily migrate UFS and other types of file system to Oracle Solaris ZFS with the provided migration tools.

Administering storage is extremely easy, because the design allows administrators to state the intent of their storage policies rather than all of the details needed to implement them. Creating a file system or performing other administrative activities is very fast—less than one second, regardless of size. There is no need to configure (or worse, reconfigure) underlying storage devices or volumes, because this is handled automatically when they are added to a pool. Oracle Solaris ZFS also allows administrators to guarantee a minimum capacity for file systems, or to set quotas to limit maximum sizes. Administrators can delegate permissions to perform Oracle Solaris ZFS administration tasks to non-privileged users at a fine-grained level, making it easy for application owners to deploy and configure Oracle Solaris ZFS quickly while still protecting overall system security and integrity.

Volume Management is A Thing of the Past

Unlike traditional file systems that require a separate volume manager, Oracle Solaris ZFS introduces the integration of volume management functions. The traditional combination of file system and volume manager maintains one-to-one mapping between the file system and its associated volumes. Oracle Solaris ZFS breaks out of this limitation with the storage pool model. When capacity is no longer required by one file system in the pool, it becomes available to others.

Reduced Costs

Oracle Solaris ZFS can reduce costs in multiple ways: decreasing the time and complexity of administrative tasks; efficiently using resources; and, eliminating licensing costs for no longer needed additional software such as volume managers. All administration tasks are performed online, resulting in zero downtime for administration. In addition, as an integrated part of Oracle Solaris, the technology does not require a separate maintenance contract. This can greatly simplify support issues; there is a single point of contact and only one maintenance contract for all software layers between the application and storage resources. And the RAID-Z feature saves money by maintaining data redundancy without the expense of a hardware RAID controller or nonvolatile RAM. Oracle Solaris ZFS supports tiered storage

configurations that include flash technology for both greater performance and cost effectiveness.

Compatibility

Applications do not need to be changed or modified to take advantage of the industry-leading capabilities of Oracle Solaris ZFS. It employs familiar POSIX interfaces, and the existing storage infrastructure—device drivers, storage fabric, and devices—works without requiring changes. For applications that prefer to operate directly on block devices, Oracle Solaris ZFS provides the *zvol* volume emulator, which delivers all of the benefits of transactional integrity and checksums as well as compatibility with existing block-based, volume manager interfaces.

Enhanced Data Services, Seamless Integration and Data Sharing

Oracle Solaris ZFS provides data encryption, deduplication, compression, infinite snapshots and clones of filesystems along with data replication across systems and sites. It is tightly integrated with higher level Oracle Solaris 11 data sharing services such as CIFS, NFS, and iSCSI. The administration of these facilities is simple and intuitive.

Conclusion

By offering data security and integrity, virtually unlimited scalability, and easy and automated manageability, Oracle Solaris ZFS simplifies storage and data management for demanding applications today—and well into the future.

Contact Us

For more information about Solaris 11, visit oracle.com/solaris or call +1.800.ORACLE1 to speak to an Oracle representative.

Oracle is committed to developing practices and products that help protect the environment

Hardware and Software, Engineered to Work Together

ORACLE